DevOps Platform Strategy

Consolidating five Azure DevOps organisations onto one governed platform

For decision: 6 items · 12-month horizon · investment and stop-list included

Internal - prepared for the CTO and the Technology Leadership Team

SECTION 01

Executive summary

The platform programme is worth finishing, but not on its current path. Re-scope it: keep migrating, replace pipeline-by-pipeline conversion with a golden path, and make the security and measurement controls real before scaling.

The situation. Five Azure DevOps organisations, roughly 1,000 repositories and ~1,500 pipelines are being consolidated onto GitHub Enterprise. Repository migration is ahead of plan, which is the good news: the hard, irreversible work is working. Pipeline conversion is behind, because converting pipelines one at a time reproduces the duplication that made the estate expensive. Security wants stronger enforcement now, several legacy products cannot use standard runners, and the metrics the programme reports are disputed - so we cannot yet prove progress to the board.

The recommendation. Re-scope rather than accelerate or pause. Acceleration buys throughput we cannot absorb while the golden path and the controls are unfinished; a pause wastes the migration lead and re-opens the estate to drift. Re-scoping moves effort from per-repository conversion to one reusable pipeline template, funds a measurement rebuild so the numbers settle disputes instead of causing them, and puts a dated exception path around the legacy runners.

What we ask of the CTO. Six decisions on page 10. Two are urgent and unblock the next quarter: approve the golden-path-first policy (replacing pipeline-by-pipeline conversion) and mandate that the DORA baseline in this brief is the single reported definition across all five organisations.

38.33	20.5 h	0.87 %	2.0 h
Deployments per project-month	Lead time for changes (median)	Change failure rate	Mean time to restore (median)
345 builds, 2 projects, 6 months	n = 149, 98.0 % link coverage	3 failed of 345 deployments	n = 3, indicative only

Read the four metrics with their limits

These are the figures this programme can currently defend: recomputed from source, reconciled row by row, with three automated quality gates passing. They are also proxies - a completed build stands in for a deployment because the source records no release event. MTTR rests on three incidents. Page 2 states exactly what is measurable today and what is not, which is what makes the numbers usable in a board paper.

- **Continue** repository migration - it is on plan and it is irreversible value.
- **Change** pipeline conversion to golden-path inheritance, not per-repository porting.
- **Enforce** non-waivable security controls on every new pipeline from day one, with a dated, owned exception path for legacy products.

What is in this brief

01 Executive summary	06 Governance and security
02 Evidence and measurement integrity	07 Twelve-month roadmap
03 Target architecture	08 The first 90 days
04 Architecture decisions	09 Risks and how they are held
05 Target operating model	10 Decisions requested from the CTO

SECTION 02

The evidence base and measurement integrity

Every figure in this brief is reproducible from the supplied dataset. Where a number is a proxy, it is labelled a proxy. A disputed metric is a governance problem, not a reporting problem.

Baseline established for the two in-scope projects, July to December 2024

Metric	Definition at this grain	Value	Evidence	Confidence
Deployment frequency	Completed builds per project-month	38.33	345 builds / 9 grain rows	High for the proxy; not a release count
Lead time for changes	Median commit-author to build-complete, hours	20.5 h	n = 149 links, 98.0 % of links resolve	High where observed; absent for 5 of 9 rows
Change failure rate	Deployment-caused incidents / deployments	0.87 %	3 incidents / 345 deployments	Directional - matched to project-month
Mean time to restore	Median incident start to resolution, hours	2.0 h	n = 3, all at 2.03 h	Indicative only

What the data can and cannot support

- **Measurable now:** deployment volume, lead time where commits link to builds, and failure rate at project-month grain. These recompute exactly and reconcile row by row.
- **Not measurable from this source:** deployments as a discrete event. There is no environment, no release marker and no deployment-to-incident link, so frequency and failure rate remain proxies. That is a source-system gap, not an analytical one.
- **Not measurable at all for one project:** the dataset contains commits for a single project. Lead time for the other cannot exist until commit-to-build linkage is captured.
- **Structurally unreachable today:** a project-month where all four metrics are observed. The only month with a deployment-caused incident has no commits. We report 4 of 9 rows as partial and 5 as minimal, and say so per row.

Disputed metric	Why the numbers disagreed	Resolution adopted
"Our deployment count is wrong"	Different teams counted pipeline runs, completed builds and releases as the same thing	Publish one definition: a completed build is the deployment proxy. Label every visual.
"Lead time is not credible"	Some reports averaged heavy-tailed data; one report included links where the build finished before the commit was authored	Median, not mean; 12 impossible associations rejected and counted, not repaired.
"Failure rate is too low"	Incidents were matched to the deployment project only where the incident carried a project path; unmapped incidents were dropped silently	Unmapped incidents are held in the dimension as explicit members, never dropped.

What becomes answerable, and when

Question the CTO will be asked	Answerable today	Answerable after the measurement rebuild
How often do we deploy?	Yes, as a completed-build proxy	Yes, as a discrete deployment event with environment
What is our failure rate?	Directionally, at project-month grain	Exactly, linked to the failing deployment
How long do we take to restore?	Indicative only - three incidents	Credible - every incident carries a deployment link
Is a given team improving?	We refuse to answer this today	Yes, with paired metrics and published sample sizes

What we propose to stop reporting

Any per-team or per-individual ranking derived from these four metrics, and any percentage that depends on a deployment event the source does not record. Both invite gaming: one rewards splitting builds, the other rewards not flagging incidents. A metric that can be improved by mislabelling is a metric that will be.

SECTION 03

Target architecture

One source of truth for change, one path to production, one measurement spine. Everything else is an exception with an owner and an end date.

Layer	What it is	How it is built	Owner
1. Source control	One GitHub Enterprise org; five ADO orgs retired	Wave migration with coexistence; history and permissions preserved	Platform
2. Change path	Golden-path pipeline inherited by every service	Reusable workflow, not copy-paste YAML; policy checked in the same PR	Platform
3. Execution	Standard hosted runners by default, tiered exceptions	GitHub-hosted; larger runners for heavy builds; self-hosted only by exception	Platform / Security
4. Telemetry	Deployment, incident and commit events captured once	Event emission from the pipeline plus incident linkage at source	Platform / SRE
5. Measurement	Bronze to Silver to Gold, then the semantic model	Ingest raw, validate and deduplicate, publish the metric grain	Data Engineering
6. Consumption	One baseline report, one definition, one owner	Power BI semantic model over Gold; no metric redefined in a dashboard	Data / DevOps

Delivery flow, end to end

Commit → pull request → automated gates (tests, SAST, dependency and licence policy, code-owner review) → golden-path pipeline → deployment event with environment → telemetry → metric grain → baseline report. Each arrow is a contract with an owner; a missing arrow is a finding, not a gap to be worked around locally.

Current state against the target, layer by layer

Layer	Today	Target at month 12	If we do nothing
Source control	5 ADO organisations, inconsistent policy	1 GitHub organisation, 1 policy set	Drift continues; audits stay manual
Change path	~1,500 pipelines, mostly copy-paste	A small number of inherited patterns	Duplication carried forward permanently
Execution	Mixed runners, unclear ownership	Hosted by default; self-hosted only by dated exception	Patching and network surface grow
Telemetry	No deployment event; incidents not linked	Every deployment emits an event; incidents linked	Metrics stay proxies; disputes persist
Measurement	Team dashboards, unreconciled	One Gold source with a reconciliation gate	Numbers cannot be defended in a board paper

Runner strategy: standard by default, exception by decision

Tier	Use	Control	Exit condition
Standard hosted	All new and migrated pipelines; the default with no approval	OIDC federation, no long-lived secrets, egress restricted	n/a - permanent baseline
Larger hosted	Builds that need more compute than the standard pool	Same controls; cost attributed to the owning product	Reviewed quarterly against spend
Self-hosted (exception)	Legacy products that cannot move to hosted runners	Registered, patched, monitored; network-isolated; owner named	Dated exception, maximum 12 months, then re-decided
Air-gapped	Only where regulation requires it	Separate estate, mirrored artefacts, no direct internet	Annual regulatory review

Non-negotiable architectural boundaries

Identity through OIDC federation with no long-lived cloud credentials in any pipeline; policy enforced as code in the same pull request that changes the pipeline; every deployment emits an event with service, environment and commit; the measurement layer reads from one Gold source. A service cannot buy its way around these with a local exception - it can only request a dated one, from the accountable owner, in writing.

SECTION 04

Architecture decisions needing endorsement

Seven decisions shape the next twelve months. Each states what we give up, because a decision without a cost is a preference.

#	Decision	Why	What we give up	Reversal cost
A1	Golden path is the only supported route to production	Converts 1,500 pipelines once instead of 1,500 times	Teams with unusual needs must request a pattern, not fork a pipeline	High after Q2
A2	Standard hosted runners are the default; self-hosted only by dated exception	Removes patching and network surface from the platform team	Some legacy products wait for a runner decision or accept a migration	Medium
A3	Security controls are enforced in the pull request, not in a quarterly audit	Prevents drift and makes the control provable on demand	Slower first merge for teams new to the gates	Low
A4	One organisation, one policy set, no per-team policy forks	Five policy sets is how the estate became inconsistent	Central priority setting; teams trade autonomy for predictability	High
A5	Metrics are published from one Gold layer, not from team dashboards	Ends the dispute about which number is true	Teams lose bespoke views that cannot be reconciled	Medium
A6	Deployment events are emitted by the pipeline and linked to incidents at source	Removes the proxy from frequency and failure rate	Short delivery work now, to stop arguing about numbers later	Low - additive
A7	Legacy products are re-platformed or explicitly retired by month 12	An exception with no end date is just a second platform	Some products must invest or accept end-of-support	High - commercial

Non-functional targets for the consolidated platform

Attribute	Target	How it is proven
Time to first green pipeline for a new service	Under 1 day from repository creation	Golden-path adoption report; median onboarding time
Change lead time (platform-owned work)	Median under 1 day for a policy-compliant change	Same metric definition as the business baseline
Availability of the change path	99.9 % in business hours; no single point of failure	Runner pool and workflow availability, measured monthly
Security control coverage	100 % of new pipelines; 95 % of migrated pipelines by day 90	Policy-as-code evaluation results, not survey answers
Metric reproducibility	Every published figure recomputable from stored facts	Automated reconciliation gate; a failure blocks publication

Alternatives considered and rejected

Alternative	Why it was attractive	Why it was rejected
Accelerate the current plan	Buys visible throughput; shows momentum	Throughput into a path we have not finished building recreates the duplication at speed
Pause the programme	Removes the friction complaints immediately	Discards the migration lead, resumes with the same problems plus a year of drift
Convert pipeline by pipeline, faster	Measurable in pipelines per week	1,500 conversions preserve 1,500 variations, and each needs its own security review
One platform team does everything	Single accountable owner; no coordination cost	Becomes a ticket queue; product teams stop owning their own reliability

The expensive mistake we are avoiding

Converting pipelines repository by repository looks like progress and is measurable in pipelines per week, which is why it is attractive. It also carries the duplication forward and leaves 1,500 slightly different security postures to keep patched. Wave 0 measures the duplicate share first: if it is as high as the sample suggests, conversion is the wrong unit of work and the plan on page 7 has to stand.

SECTION 05

Target operating model

One platform team owns the product; federated squads consume it. Accountability is named per area, not distributed across a committee.

Area	Accountable	Responsible for delivery	Consulted
Platform product (golden path, runners, templates)	Head of DevOps & Analytics	Platform engineering	Security, product engineering leads
Security controls and exceptions	Chief Information Security Officer	Security engineering	Platform, legal, product
Measurement and reporting integrity	Head of DevOps & Analytics	Data engineering	Finance, product leadership
Repository and pipeline migration waves	Programme lead (Engineering)	Migration squad	Platform, product owners
Service reliability in production	Product engineering leads	Owning squad on call	SRE, platform
Legacy re-platform or retirement decisions	CTO	Product owners	Finance, platform, commercial

Service tiers: commitments we can keep

Tier	What sits here	Platform commitment	Support
T0 - Revenue critical	Customer-facing transaction paths	Golden path with change freeze windows; 99.9 % change-path availability	24x7 on call, 15-minute platform response
T1 - Business critical	Internal systems with revenue dependency	Golden path; 99.5 % change-path availability	Extended hours, 1-hour response
T2 - Standard	Everything else	Golden path, best effort within business hours	Business hours, planned work
Exception	Legacy products that cannot meet T0-T2 controls	Documented workaround with an owner and an expiry date	Per exception agreement, reviewed monthly

Decision rights, so the programme stops escalating

Decision	Decides	Advises	Escalates to
Approve a new golden-path pattern	Head of DevOps & Analytics	Platform, security	CTO if it changes a control
Grant a security exception	CISO	Platform, product owner	CTO and CISO jointly, if over 90 days
Change a published metric definition	Head of DevOps & Analytics	Finance, data	CTO
Pause a migration wave	Programme lead	Platform, product	CTO
Re-platform or retire a legacy product	CTO	Product, finance	Executive committee

Capability gaps between today and this model

Capability	Gap today	How it closes	By
Platform product management	No named owner; backlog driven by tickets	Appoint a product owner with authority to say no to off-catalogue work	Month 1
Policy as code	Controls described in documents, checked quarterly	Encode each control in the change path; audit becomes sampling	Month 2
Deployment event emission	Pipelines produce no deployment record	Add event emission to the golden path; backfill is not required	Month 3
Metric stewardship	Definitions differ per team	One specification with an owner, a version and a reconciliation gate	Month 2

The internal developer platform is a product, not a project

Three personas consume it: the service team that needs a compliant pipeline in a day; the security engineer who needs controls provable on demand; and the engineering leader who needs a number they can defend. Catalogue v1 is deliberately small - service scaffold, golden-path pipeline, environment provisioning, deployment event emission, metric self-service. Anything outside the catalogue is a request with a business case, which is how the platform team avoids becoming a ticket queue.

SECTION 06

Governance and security

Controls are enforced where change happens - the pull request - so compliance becomes a by-product of delivery rather than a quarterly scramble.

Control	Requirement	Enforced by	Evidence produced
Branch protection	No direct push to main; review by code owner; status checks green	Repository ruleset, evaluated on push	Ruleset evaluation log per repository
Credential hygiene	No long-lived cloud secrets in any pipeline; OIDC federation only	Secret scanning plus pipeline policy	Zero-secret attestation per pipeline
Dependency risk	Known-vulnerable and licence-prohibited packages block the merge	Dependency review plus audit gate	Blocked-merge report, per pull request
Static analysis	Security findings above threshold block release	Code scanning in the PR	Finding disposition, tracked to closure
Pipeline policy	A pipeline may only inherit approved patterns	Policy-as-code check in the PR	Policy verdict attached to the commit
Data handling	Personal data never leaves Silver; Gold is publishable	Pipeline contract plus review	Column-level lineage and drop log
Metric integrity	A published figure must recompute from stored facts	Automated reconciliation gate	Gate result with sample counts
Audit trail	Who changed what, when, with which approval	Immutable audit log, retained per policy	Extract on demand for audit

Exception process: fast, dated and owned

Step	Timebox	Who	Outcome
1. Request with a reason and an end date	Same day	Product owner	Ticket with the failing control named
2. Risk assessment against the control	2 working days	Security engineering	Accept, mitigate or decline, in writing
3. Time-boxed approval, maximum 90 days	1 working day	CISO (or CTO above 90)	Exception registered, visible on a dashboard
4. Review before expiry	Monthly	CISO and platform	Renew once with justification, or close

Metric governance: make gaming harder than doing the work

- **One definition, one owner.** A metric cannot be redefined inside a dashboard. Any change to formula, grain or exclusion is a version bump, a pipeline change and a release note.
- **Pairs, not singles.** Report frequency with failure rate; report lead time with sample size. A single number can always be improved by degrading its partner.
- **Sample sizes are published, never hidden.** A median over three observations is labelled indicative, in the report, not in a footnote.
- **Blank is not zero.** A metric with no observation in a period stays blank. Zero-filling is the most common way a reliability metric becomes a fiction.

> **The governance question the CTO will be asked**
> "Who decided this control applies to us?" The answer must be a named role with a written rationale, not a committee. That is why decision rights are on page 5 and every control above has a single enforcing mechanism: if two systems enforce the same rule, neither is accountable when it fails.

SECTION 07

Twelve-month roadmap

Waves are sequenced by dependency and by evidence, not by enthusiasm. Each wave has an exit criterion that can fail.

Wave	Window	Scope	Exit criterion (falsifiable)
Wave 0	Days 1-30	Baseline the estate: duplicate-pipeline analysis, runner inventory, control gap assessment, metric reconciliation	Duplicate share quantified; every legacy runner has an owner and a date
Wave 1	Months 2-3	Golden path v1 for the two highest-volume service patterns; security controls live on all new pipelines	10 services on the golden path; 100 % of new pipelines pass every control
Wave 2	Months 4-6	Migrate the high-volume, low-exception repositories; retire the first ADO organisation	First org decommissioned; no regression in the four DORA metrics
Wave 3	Months 7-9	Remaining standard repositories; legacy products re-platformed or given a dated decision	80 % of repositories migrated; no open-ended exceptions remain
Wave 4	Months 10-12	Final consolidation, decommission the last ADO organisation, hand over to run	One organisation in use; cost reduction realised; run model signed off

Milestones the leadership team will hold us to

ID	Milestone	Due	Owner	Status today
M1	Security controls enforced on 100 % of new pipelines	Month 1	CISO / Platform	Not started - needs the decision on page 10
M2	Golden path adopted by 10 services	Month 3	Platform	Design agreed
M3	Measurement baseline published from one Gold source	Month 2	Data engineering	Achieved for the in-scope projects
M4	First ADO organisation decommissioned	Month 6	Programme lead	Planned
M5	80 % of repositories migrated	Month 9	Programme lead	Migration ahead of plan
M6	Single organisation, cost reduction realised	Month 12	CTO	Planned

Decommissioning: an organisation is only retired when it is empty

- No repository has pushed to the old organisation for 30 consecutive days.
- No pipeline in the old organisation has run for 30 consecutive days, other than read-only audit jobs.
- All history, work items and permissions are either migrated or explicitly archived with a named owner and a retention decision.
- Billing, identity and audit integrations are cut over and the security team has signed off.

What each wave gives the business

Wave	Business outcome	How it is measured
Wave 0	A defensible number and an honest inventory	Duplicate share quantified; every legacy runner owned and dated
Wave 1	New services reach production in a day, compliantly	Median time to first green pipeline; control pass rate at 100 %
Wave 2	First platform retired: cost and audit surface fall	Organisation decommissioned; run cost variance
Wave 3	Majority of the estate on one governed path	Repositories migrated; open-ended exceptions at zero
Wave 4	Single platform, predictable cost, defensible metrics	One organisation in use; DORA reported monthly without dispute

The dependency that decides the year

Wave 1 cannot start at scale until the golden path exists, and the golden path cannot be trusted until the security controls are enforced inside it. That is why M1 sits in month one and why the stop-list on page 10 removes work from the same team. Sequencing these three in the wrong order is the most likely way this programme loses a quarter.

SECTION 08

The first 90 days

Enough detail to start on Monday, and a day-90 gate that can stop the programme if the evidence is not there.

Window	Platform	Security	Measurement	Migration
Days 1-15	Golden-path design review; duplicate-pipeline analysis	Control catalogue agreed; policy-as-code prototypes	Baseline re-published from Gold; definitions signed off	Wave 1 repository shortlist with owners
Days 16-30	Golden path v1 built; two pilot services onboarded	Controls live on new pipelines in enforce mode	Reconciliation gate wired into publication	Legacy runner inventory with expiry proposals
Days 31-60	Golden path adopted by 6 services; template versioning	First exception register published	Deployment-event schema agreed with platform	First repositories migrated on the golden path
Days 61-90	10 services on the golden path; catalogue v1 published	Control coverage measured, not asserted	DORA baseline reported to leadership monthly	First ADO organisation ready for decommission review

Day-90 gate: what must be true to continue at this investment

Gate	Threshold	If missed
Golden-path adoption	At least 10 services running on it	Re-scope to one pattern; do not widen
Security control coverage	100 % of new pipelines; measured, not asserted	Stop onboarding new services until closed
Metric reproducibility	Every published figure recomputes; gate green	Withdraw the disputed figure from reporting
Migration progress	At least one organisation decommission-ready	Re-plan the wave model with the CTO
Developer friction	No increase in median change lead time for onboarded teams	Fix the golden path before adding services

Staffing assumption behind the plan

- **Platform engineering:** 4 engineers, with at least 2 able to own the golden path as a product rather than a script.
- **Security engineering:** 1 dedicated engineer for policy-as-code and the exception register - not a fraction of someone's week.
- **Data engineering:** 1 engineer for the measurement spine and the reconciliation gate.
- **Migration squad:** 2 engineers plus product-side owners per wave.
- **Product management:** a named owner for the platform product, accountable for the catalogue and for saying no.

What we deliberately will not do in the first 90 days

- Migrate anything that cannot run on the golden path. Migrating an exception just moves it.
- Backfill historical deployment events. The baseline already states what is measurable; inventing history would undermine it.
- Publish per-team rankings. The definitions cannot yet support them, and the incentive is wrong.
- Retire a legacy runner before its product has a funded route to hosted execution.
- Take on new repositories outside the agreed wave list, however small the request looks.

> **Why 90 days and not 12 months**
> A year-long plan with no intermediate gate is a budget request, not a plan. The day-90 gate tests the two assumptions everything else rests on: that teams will adopt a golden path, and that the controls can be enforced without stopping delivery. If either is false, the right response is to change the approach while it is still cheap.

SECTION 09

Risks and how they are held

Each risk has an owner, an early warning that can be seen before the impact, and a decision it would force.

#	Risk	Likelihood	Impact	Mitigation and early warning
R1	Golden path is not adopted; teams keep forking pipelines	Medium	High	Product owner with authority to say no; adoption measured weekly; warning = fork count rising
R2	Security enforcement slows delivery and creates shadow pipelines	Medium	High	Exception path with a 3-day turnaround; warning = exception volume and off-platform runs
R3	Legacy products block decommissioning indefinitely	High	Medium	Dated exceptions, maximum 12 months, re-decided by the CTO; warning = exceptions past expiry
R4	Metric definitions stay disputed, so progress cannot be proven	Medium	High	One Gold source, reconciliation gate, published definitions; warning = dashboard divergence reports
R5	Migration consumes the platform team and the golden path slips	High	High	Stop-list on page 10; wave capacity reserved for platform; warning = golden-path backlog age
R6	Key-person dependency in platform and data engineering	Medium	Medium	Two owners per component; documented runbooks; warning = single-owner components in review
R7	Cost of consolidation exceeds the saving in year one	Medium	Medium	Cost tracked per wave against baseline; warning = run-cost variance above 10 %
R8	Developer friction erodes trust in the programme	High	Medium	Lead time measured for onboarded teams; quarterly developer survey; warning = lead time regression

What would make us change the recommendation

Signal	Response	Decided by
Golden path rejected by more than a third of pilot teams	Pause expansion; run a discovery on why, then re-scope the product	CTO with platform
Security exception volume grows week on week	Treat as a control design failure; fix the control, not the teams	CISO
Metric dispute unresolved after the reconciliation gate	Publish only what reconciles; withdraw the rest from leadership reporting	Head of DevOps
Migration stalls for two consecutive waves	Re-plan the wave model; consider consolidating to three organisations first	CTO

Risk appetite, stated plainly

Category	Appetite	Consequence for the programme
Security	Very low - no new pipeline without enforced controls	We accept slower onboarding rather than a known-unsafe path
Delivery pace	Moderate - waves slip before controls slip	A missed wave date is a scheduling change; a missed control is an incident
Legacy exceptions	Low, and strictly time-bound	An exception without an expiry date is treated as a failed decision, not a state
Metric precision	High - we would rather publish blank than wrong	Any figure that cannot be recomputed is withdrawn from leadership reporting

The risk that matters most

R5 is the one that decides the year: the same small team is being asked to migrate an estate, build a platform and fix the measurement. Something must come off that team's plate, and the stop-list on page 10 is the specific proposal. Left unaddressed, the predictable outcome is a half-migrated estate with a golden path nobody finished and metrics still in dispute - the most expensive of the available endings.

SECTION 10

Decisions requested from the CTO

Six decisions. Two are needed this month to unblock Wave 1; the rest shape the year and can be taken with the quarterly plan.

#	Decision	Options considered	Recommendation	Needed by
D1	Adopt golden-path-first conversion	Keep converting per repository / adopt inheritance / pause conversion	**Adopt inheritance**, with conversion frozen except for Wave 1 pilots	Month 1
D2	Mandate one DORA definition across all organisations	Let each org keep its own / adopt the measurement baseline in this brief	**Adopt the baseline**, published from one Gold source, with the stated limits	Month 1
D3	Enforce non-waivable security controls on all new pipelines	Enforce immediately / enforce with exceptions / educate first	**Enforce with a dated exception path**, 3-day turnaround	Month 1
D4	Fund the legacy runner exception path	Force all products to hosted runners / fund self-hosted isolation / extend timelines	**Fund isolation** for named products, maximum 12 months	Month 2
D5	Re-platform or retire unmodernisable legacy products	Defer / invest / retire	**Decide per product by month 6**, with a commercial owner named	Month 6
D6	Approve the stop-list	Continue everything / stop the three items on the right	**Approve the stop-list** to create platform capacity	Month 1

Investment shape, and what it displaces

Item	Direction	Basis
Platform team (4 engineers, 12 months)	Funded from the migration programme	Reallocated from per-repository conversion effort, not new headcount
Security engineering (1 engineer)	New, ring-fenced	Policy-as-code and the exception register; not a fraction of an existing role
Data engineering (1 engineer)	Part-funded from reporting consolidation	Retires several team-level dashboards and their maintenance
Legacy runner isolation	Time-boxed, product-attributed	Cost sits with the product that needs the exception, not with the platform
Tooling and licences	Neutral to down	Consolidation reduces duplicate tooling across five organisations

What we stop, to make room

- **Stop** per-repository pipeline porting outside Wave 1 pilots. This is the single largest source of avoidable work and of future duplication.
- **Stop** publishing team-level DORA dashboards that cannot be reconciled to Gold.
- **Stop** quarterly security audits as the primary control, replaced by enforcement in the pull request plus a sampled audit.

If the four decisions are approved this month

When	What happens	Evidence the CTO will see
Week 1	Golden-path-first policy published; conversion frozen outside pilots	Policy note and the amended wave list
Week 2-4	Controls turned on in enforce mode for new pipelines; exception register live	Coverage report and the first exceptions, with owners and dates
Month 2	DORA baseline published from one Gold source with its stated limits	The report in this repository and the reconciliation gate result
Month 3	Day-90 gate review against the five thresholds on page 8	Adoption, coverage, reproducibility, migration and friction evidence

The ask, in one sentence

Approve D1, D2, D3 and D6 this month so Wave 1 can start on a golden path with enforceable controls and one agreed number, and accept the stop-list that funds the team to do it. Everything else in this brief follows from those four decisions.